

# Variational Autoencoders

Elia Bronzo, Shoichi Yip

05/03/2024

- 1 VAEs from a probabilistic prospective
  - Learning probabilistic models
  - Latent variable models
  - Learning latent variable models with Expectation-Maximization
  - Learning latent variable models with Variational Auto-Encoders
- 2 VAEs from the perspective of neural networks
  - Autoencoder Neural Networks
  - Variational Autoencoder Neural Networks
  - Why use Neural Networks?
  - Finding the loss function
  - Computing the  $\mathcal{L}$
  - The reparametrization trick
  - Loss function and NN structure

# Learning probabilistic models

Let's start by fixing some notation:

- $\mathbf{x}^{(i)} \in \mathcal{D}$  that we assume to be  $N$  i.i.d. from  $p^*(\mathbf{x})$   $i = 1, \dots, N$
- $p^*(\mathbf{x})$  is the *True* generative process we want to learn

Let's start considering a **fully observed model**

$$p_{\theta}(\mathbf{x}) \approx p^*(\mathbf{x})$$

## Maximum Likelihood (ML)

$$\begin{aligned}\theta_{ML} &= \arg \max_{\theta} p_{\theta}(\mathcal{D}) \\ &= \arg \max_{\theta} \sum_{i=1}^N p_{\theta}(\mathbf{x}^{(i)})\end{aligned}$$

## Maximum a posteriori (MAP)

$$\begin{aligned}\theta_{MAP} &= \arg \max_{\theta} p(\theta|\mathcal{D}) \\ &= \arg \max_{\theta} \frac{p_{\theta}(\mathcal{D})p(\theta)}{p(\mathcal{D})}\end{aligned}$$

# Why not using DNNs?

$p_\theta \leftarrow$  Deep Neural Networks (DNN)

- high representation capacity
- easily trained with SGD

SGD:

$$\frac{1}{N_{\mathcal{D}}} \nabla_{\theta} \log p_{\theta}(\mathcal{D}) \quad \underbrace{\approx}_{\text{RHS unbiased estimator of LHS}} \quad \underbrace{\frac{1}{N_{\mathcal{M}}} \nabla_{\theta} \log p_{\theta}(\mathcal{M})}_{\text{Gradient over minibatch}}$$

# Latent variables models

$p_{\theta}(\mathbf{x}, \mathbf{z})$  where  $\mathbf{z}$  latent variables

$$p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{x}, \mathbf{z}) d\mathbf{z} = \int \underbrace{p_{\theta}(\mathbf{z})}_{\text{Prior}} p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z} \quad \text{Marginal likelihood}$$

If  $p_{\theta}(\mathbf{x}|\mathbf{z})$  is a DNN  $\rightarrow$  **Deep Latent Variable model (DLVM)**

Why?

- high representation capacity:  $p_{\theta}(\mathbf{x}) \rightarrow$  (almost) arbitrary dependencies even with simple (ex. Gaussian) priors
- easily trained with SGD

**VAE** is used to learn this kind of models

# Learning latent variable models

Learning a latent variable model means optimize the **marginal log-likelihood**:

$$l(\theta) = \sum_{i=1}^N \log p_{\theta}(\mathbf{x}^{(i)}) = \sum_{i=1}^N \int p_{\theta}(\mathbf{z}) p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) d\mathbf{z}$$

The *classical* way of optimizing in this problems is **Expectation-Maximization**

We will first focus on how to optimize  $l(\theta)$  in case of a dataset composed by only 1 point ( $\mathbf{x}^{(1)}$ ) ( $N=1$ )

# Evidence Lower Bound (ELBO)

$$\begin{aligned}
 \log p_{\theta}(\mathbf{x}^{(1)}) &= \log \int d\mathbf{z} p_{\theta}(\mathbf{x}^{(1)}, \mathbf{z}) \\
 &= \log \underbrace{\int d\mathbf{z} q(\mathbf{z}) \frac{p_{\theta}(\mathbf{x}^{(1)}, \mathbf{z})}{q(\mathbf{z})}}_{\mathbb{E}_q} \\
 &\underbrace{\geq}_{\text{Jensen's}} \underbrace{\int d\mathbf{z} q(\mathbf{z}) \log \frac{p_{\theta}(\mathbf{x}^{(1)}, \mathbf{z})}{q(\mathbf{z})}}_{\text{ELBO}}
 \end{aligned}$$

This is true **for any**  $q(\mathbf{z})$

**Jensen's inequality**  $f$   
convex:

$$\mathbb{E}[f(X)] \geq f(\mathbb{E}[X])$$

- tight bound  $\iff f$  strictly convex
- $\log(x)$  concave

## More on the ELBO

$$\begin{aligned}\log p_{\theta}(\mathbf{x}^{(1)}) &= \mathbb{E}_q \left[ \log p_{\theta}(\mathbf{x}^{(1)}) \right] \\ &= \mathbb{E}_q \left[ \log \left[ \frac{p_{\theta}(\mathbf{x}^{(1)}, \mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x}^{(1)})} \right] \right] \\ &= \mathbb{E}_q \left[ \log \left[ \frac{p_{\theta}(\mathbf{x}^{(1)}, \mathbf{z})}{q(\mathbf{z})} \frac{q(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x}^{(1)})} \right] \right] \\ &= \mathbb{E}_q \left[ \log \left[ \frac{p_{\theta}(\mathbf{x}^{(1)}, \mathbf{z})}{q(\mathbf{z})} \right] \right] + \mathbb{E}_q \left[ \log \left[ \frac{q(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x}^{(1)})} \right] \right] \\ &= \text{ELBO}(\mathbf{x}^{(1)}; q, \theta) + \mathcal{D}_{\text{KL}} \left[ q(\mathbf{z}) \parallel p_{\theta}(\mathbf{z}|\mathbf{x}) \right]\end{aligned}$$

When does the ELBO is **tight**? (equality)

$$\text{ELBO}(\mathbf{x}^{(1)}; q, \theta) = \log p_{\theta}(\mathbf{x}^{(1)}) \iff q(\mathbf{z}) = p_{\theta}(\mathbf{z}|\mathbf{x}^{(1)})$$



# Expectation-Maximization Algorithm

- **Expectation** step: at fixed  $\theta^{(t)}$ ,  $q(\mathbf{z}) \leftarrow p_{\theta^{(t)}}(\mathbf{z}|\mathbf{x}^{(1)})$
- **Maximization** step: at fixed  $q$ ,  $\theta^{(t+1)} \leftarrow \arg \max_{\theta} \text{ELBO}(\mathbf{x}^{(1)}; q, \theta)$

What about a dataset with more than 1 point?  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$

$$l(\theta) = \sum_{i=1}^N \log p_{\theta}(\mathbf{x}^{(i)}) \geq \sum_{i=1}^N \text{ELBO}(\mathbf{x}^{(i)}; q_i, \theta)$$

Repeat until convergence {

(E-step) **for each i:**

$$q_i(\mathbf{z}) \leftarrow p_{\theta^{(t)}}(\mathbf{z}|\mathbf{x}^{(i)})$$

(M-step)

$$\theta^{(t+1)} \leftarrow \arg \max_{\theta} \sum_{i=1}^N \text{ELBO}(\mathbf{x}^{(i)}; q, \theta) \}$$

# DLVM, intractabilities and problems with EM

$$l(\theta) = \sum_{i=1}^N \log p_{\theta}(\mathbf{x}^{(i)}) = \sum_{i=1}^N \underbrace{\int p_{\theta}(\mathbf{z}) p_{\theta}(\mathbf{x}^{(i)} | \mathbf{z}) d\mathbf{z}}_{\text{intractable for a DNN}}$$

Intractable means that we don't have neither an *analytical solution* nor an *efficient estimator* of the integral in case of a DNN

Intractable  $p_{\theta}(\mathbf{x}) \iff$  Intractable **posterior**  $p_{\theta}(\mathbf{z} | \mathbf{x}) = \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{p_{\theta}(\mathbf{x})} = \frac{p_{\theta}(\mathbf{z}) p_{\theta}(\mathbf{x} | \mathbf{z})}{p_{\theta}(\mathbf{x})}$

E-step: evaluation of an intractable posterior with MCMC **for each data-point**  $\rightarrow$  computationally expensive

Solutions?  $\rightarrow$  **VAE**

# Variational Auto-Encoders (VAE)

Model for the posterior  $q_{\phi}(\mathbf{z}|\mathbf{x}) \rightarrow$  another DNN

# Latent representations





Figure: Chess memory experiment and structure of an autoencoder (from HOML).

# Undercompleteness



**Figure:** The typical structure of an autoencoder and of an MLP.

# What is the autoencoder learning?

The autoencoder learns the **identity**. We expect the output to resemble to the input.

But it is characterised by the **information bottleneck** that yields **undercompleteness**.

Depending on the characteristics of the bottleneck, we get different autoencoders.

# Stacked autoencoder




Figure: Structure of a stacked autoencoder (from HOML).

# Reconstruction for the stacked autoencoder



**Figure:** Reconstruction of images on a stacked autoencoder (from HOML).



# Denoising autoencoder

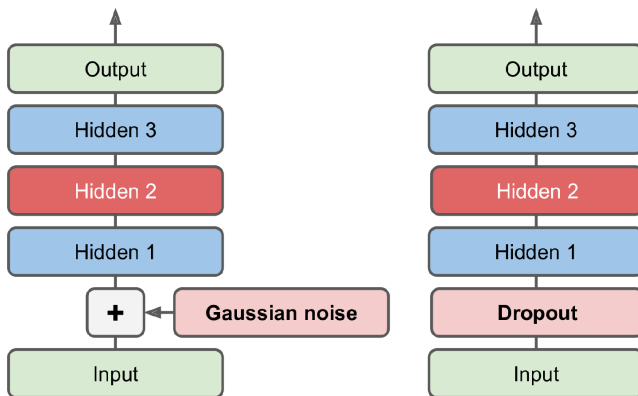


Figure: Structure of a denoising autoencoder (from HOML).

# Reconstruction of noiseless images



**Figure:** Reconstruction of noiseless images on a denoising autoencoder (from HOML).

# Pros and Cons

What autoencoders **do**:

- learn **latent** representation
- **semisupervised learning**
- **pretraining** using greedy layerwise approach
- **denoising** images

# Pros and Cons

What autoencoders **do not do**:

- have a natural **interpolation** scheme
- have **generative** capacity

# Variational autoencoder

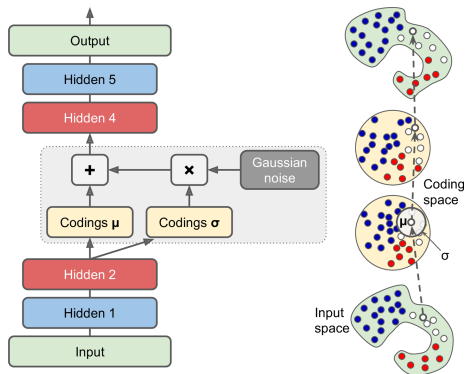


Figure: Structure of a variational autoencoder (from HOML).

# The decoder

We have said that our model is **generative**.

The generative role of the variational autoencoder is played by the **decoder**. The role of the decoder is to generate new samples that belong to the input space  $\mathcal{X}$  starting from a given vector  $\mathbf{z}$  that belongs to the space of latent variables.

The role of the decoder network is to learn the **nonlinear mapping**  $p_{\theta}(\mathbf{x}|\mathbf{z})$  that associates a value of  $\mathbf{x}$  given a point in the latent space  $\mathbf{z}$ .<sup>1</sup>

---

<sup>1</sup>This is justified by the **universal approximation theory** that states that deep ANNs are able to approximate any function.

# The decoder

We can make the assumption that the  $p(\mathbf{x}|\mathbf{z})$  is a multivariate Gaussian distribution, for example in the case of MNIST datasets.

$$p_{\theta}(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{x}|\mu_{\theta}(\mathbf{z}), \sigma_{\theta}^2(\mathbf{z}))$$

Inputs in that case are, in fact, vectors of real valued variables (i.e. pixel brightness).

## Back to the probabilistic picture

Remember that we are trying to maximise the **likelihood**  $p(x_i)$  for each  $i$ -th data point.

Given a new probability distribution  $q_i(\mathbf{z})$  we have

$$\begin{aligned} p(x_i) &= \int p(x_i|\mathbf{z})p(\mathbf{z}) \, d\mathbf{z} \\ &= \int p(x_i|\mathbf{z})p(\mathbf{z}) \frac{q_i(\mathbf{z})}{q_i(\mathbf{z})} \, d\mathbf{z} \\ &= \left\langle \frac{p(x_i|\mathbf{z})p(\mathbf{z})}{q_i(\mathbf{z})} \right\rangle_{\mathbf{z} \sim q_i(\mathbf{z})} \end{aligned}$$



## Back to the probabilistic picture

Per Jensen, we have that the **log likelihood** has the following bound

$$\begin{aligned}\log(p(x_i)) &\geq \left\langle \log \frac{p(x_i|\mathbf{z})p(\mathbf{z})}{q_i(\mathbf{z})} \right\rangle_{\mathbf{z} \sim q_i(\mathbf{z})} \\ &\geq \langle \log p(x_i|\mathbf{z})(p(\mathbf{z})) \rangle_{\mathbf{z} \sim q_i(\mathbf{z})} - H[q_i(\mathbf{z})]\end{aligned}$$

The RHS is what we have previously called ELBO (Evidence Lower Bound). We choose to maximise the ELBO instead of the log likelihood. But we are maximising against a **lot** of parameters: we have to optimise against  $\theta$  and against each  $i$ -th  $q_i(\mathbf{z})$ !

Can we reduce the number of parameters to optimise against?

# Introducing: the encoder network

What if we add another network, but this time to perform the **inference**?  
This means that we **choose** the  $q_i(z)$  to be

$$q_i(z) = q_\phi(\mathbf{z}|\mathbf{x})$$

where we call  $q_\phi(\mathbf{z}|\mathbf{x})$  the **encoder** and  $\phi$  is the set of parameters of the variational inference.

# Deep and amortised

Let us recap. In one go we

- ① chose to perform the **deep inference**, which implies that we will train a neural network to **predict** the variational parameters  $\phi$  instead of getting them by directly optimising the ELBO.
- ② reduced the number of parameters that we have to optimise against since we do not have anymore  $N$  approximate posteriors  $q_i(z)$ 's, but one parameter set  $\phi$  shared by all the data points. Since the task of inference is amortised across the entire dataset, this technique is called **amortised variational inference**.

# The Kullback-Leibler divergence

We are going to compute the **Kullback-Leibler divergence** for a particular case, in which the real and approximate posteriors have take the **multinomial Gaussian** form.

$$\mathcal{D}_{\text{KL}}[q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z})] = \int d\mathbf{z} q_{\phi}(\mathbf{z}) [\log p_{\theta}(\mathbf{z}) - \log q_{\phi}(\mathbf{z})] \quad (1)$$

This assumption can be made, for example, in the case of the MNIST dataset, since pixels are labeled with values that range from 0 to 255 depending on the intensity of the pixel. If pixels were binary variables, a Bernoulli distribution assumption would have been more suited.

# First term of $D_{\text{KL}}$

$$\begin{aligned} \int q_{\phi}(\mathbf{z}) \log p(\mathbf{z}) \, d\mathbf{z} &= \\ &= \int d\mathbf{z} \frac{1}{\sqrt{\det(\Sigma)} \sqrt{2\pi}^L} \prod_{l=1}^L \exp\left\{-\frac{(z_l - \mu_l)^2}{2\sigma_l^2}\right\} \log \frac{1}{\sqrt{2\pi}^L} e^{-\frac{|\mathbf{z}|^2}{2}} \end{aligned}$$

# First term of $D_{\text{KL}}$

$$\begin{aligned} \int q_{\phi}(\mathbf{z}) \log p(\mathbf{z}) \, d\mathbf{z} &= \\ &= \int d\mathbf{z} \frac{1}{\sqrt{\det(\Sigma)} \sqrt{2\pi}^L} \prod_{l=1}^L \exp\left\{-\frac{(z_l - \mu_l)^2}{2\sigma_l^2}\right\} \left[-\frac{L}{2} \log 2\pi - \frac{|\mathbf{z}|^2}{2}\right] \\ &= -\frac{L}{2} \log 2\pi - \frac{1}{2} \sum_{l=1}^L (\sigma_l^2 + \mu_l^2) \end{aligned}$$

## Second term of $D_{\text{KL}}$

$$\begin{aligned} & \int q_{\phi}(\mathbf{z}) \log q_{\phi}(\mathbf{z}) \, d\mathbf{z} \\ &= \int d\mathbf{z} \frac{1}{\sqrt{\det(\Sigma)} \sqrt{2\pi}^L} \prod_{l=1}^L \exp\left\{-\frac{(z_l - \mu_l)^2}{2\sigma_l^2}\right\} \\ & \quad \log \frac{1}{\sqrt{\det(\Sigma)} \sqrt{2\pi}^L} \exp\left\{-\frac{(z_l - \mu_l)^2}{2\sigma_l^2}\right\} \end{aligned}$$

## Second term of $D_{\text{KL}}$

$$\begin{aligned} & \int q_{\phi}(\mathbf{z}) \log q_{\phi}(\mathbf{z}) \, d\mathbf{z} \\ &= \int d\mathbf{z} \frac{1}{\sqrt{\det(\Sigma)} \sqrt{2\pi}^L} \prod_{l=1}^L \exp \left\{ -\frac{(z_l - \mu_l)^2}{2\sigma_l^2} \right\} \\ & \quad \left[ -\frac{L}{2} \log 2\pi - \log \sigma_l^2 - \frac{(z_l - \mu_l)^2}{2\sigma_l^2} \right] \\ &= -\frac{L}{2} \log 2\pi - \frac{1}{2} \sum_{l=1}^L (1 + \log \sigma_l^2) \end{aligned}$$



# Final expression for the Kullback-Leibler divergence

## KL divergence for the Gaussian case

The KL divergence in the Gaussian case is given by

$$\mathcal{D}_{\text{KL}}[q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p_{\theta}(\mathbf{z})] = \frac{1}{2} \sum_{l=1}^L [1 + \log \sigma_l^2 - \sigma_l^2 - \mu_l^2] \quad (2)$$

We have hence obtained the **tightness** term of our loss function given by the **ELBO**.

# The loss function

## Loss function for the Gaussian case

The loss function for the Gaussian case is given by

$$\mathcal{L}_{\theta,\phi}(\mathbf{x}^{(i)}) = (\mathbf{x}^{(i)} - \hat{\mathbf{x}}^{(i)})^2 - \frac{1}{2} \sum_{l=1}^L [1 + \log \sigma_l^2 - \sigma_l^2 - \mu_l^2] \quad (3)$$

In this simplified case, hence, we can get an closed explicit form for the loss function. It can be directly implemented in the neural network.

# A nondeterministic component

Remember the structure of the variational autoencoder.




Figure: Structure of a variational autoencoder (from HOML).

The part in grey is a **nondeterministic** component, since it has a noise term. This makes backpropagation **unfeasible**! How shall we deal with it?

# Framing the problem

Formally, what bothers us is that when we are backpropagating weight updates through the neural network, we want to be able to compute

$$\theta^{\text{new}} = \theta^{\text{old}} + \eta \nabla_{\theta} \text{ELBO}_{\theta, \phi}$$

$$\phi^{\text{new}} = \phi^{\text{old}} + \eta \nabla_{\phi} \text{ELBO}_{\theta, \phi}$$

where

$$\text{ELBO}_{\theta, \phi} = \langle \log p(x_i | \mathbf{z})(p(\mathbf{z})) \rangle_{\mathbf{z} \sim q_{\phi}(\mathbf{z} | \mathbf{x})} - H[q_{\phi}(\mathbf{z} | \mathbf{x})]$$

First derivation is **OK!**

Second derivation is **NOT OK!** Expectation value and gradient do not commute in this case. What can we do?

# Let it commute!

Let us **make it commute!** In fact, in particularly simple settings such as when the approximate posteriors  $q_\phi(\mathbf{z}|\mathbf{x})$  have Gaussian form

$$q_\phi(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}(\mathbf{x}), \boldsymbol{\sigma}(\mathbf{x}))$$

the  $\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})$  can be rewritten as

$$\mathbf{z} = \boldsymbol{\mu}(\mathbf{x}) + \boldsymbol{\epsilon} \cdot \boldsymbol{\sigma}(\mathbf{x})$$

where  $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{1})$ .

# The reparametrization trick

Instead of sampling from the approximate variational posterior we sample  $\epsilon$  from a base noise distribution and apply a **deterministic transformation**.

This amounts to replacing  $\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})$  with  $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{1})$  and defining the new deterministic  $\mathbf{z} = \boldsymbol{\mu} + \epsilon\sigma$ . It is called **reparametrization trick**

$$\begin{aligned}\nabla_\phi \text{ELBO}_{\theta,\phi} &= \nabla_\phi \langle \log p(\mathbf{x}, \boldsymbol{\mu} + \epsilon\sigma) \rangle_{\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{1})} - \nabla_\phi H[q_\phi(\mathbf{z}|\mathbf{x})] \\ &= \langle \nabla_\phi \log p(\mathbf{x}, \boldsymbol{\mu} + \epsilon\sigma) \rangle_{\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{1})} - \nabla_\phi H[q_\phi(\mathbf{z}|\mathbf{x})]\end{aligned}$$

# Structure of our network




Figure: Structure of our variational autoencoder

# Let's get to work

The Colab notebook is available at the following URL  
[https://bit.ly/pcsm1\\_vae](https://bit.ly/pcsm1_vae)